Compare the quality of forecasting models for value at risk

Author:
Hafees Adebayo YUSUFF

Supervisor:
Prof. Ralf KORN

July 23, 2021

TECHNISCHE UNIVERSITÄT
KAISERSLAUTERN

This thesis is written as a requirement for the completion of my degree of Master of Science at Technical University of Kaiserslautern.

Contents

Chapter 1

Introduction

1.1 Motivation

Basel I (Basel Accord) is the agreement reached in 1988 in Basel (Switzerland) by the Basel Committee on Bank Supervision (BCBS), involving the chairmen of the central banks of some European countries and the United States of America. This accord provides recommendations on banking regulations with regard to credit, market and operational risks. It aims to ensure that financial institutions hold enough capital on account to meet obligations and absorb unexpected losses.

For a financial institution measuring the risk it faces is an essential task. In the specific case of market risk, a possible method of measurement is the evaluation of losses likely to be incurred when the price of the portfolio assets falls. This is what Value at Risk (VaR) does.

Value at Risk(VaR) is the most common way of measuring market risk. It determines the greatest possible loss, assuming an α significance level under a normal market condition at a set time period.

Many VaR estimation methods have been developed in order to reduced uncertainty. It is however of interest to compare these method and determine the prevalence of one VaR estimation approach over others.

1.2 Literature review

The first papers involving the comparison of VaR methodologies, such as those by Beder (1995, 1996), Hendricks (1996), and Pritsker (1997), reported that the Historical Simulation performed at least as well as the methodologies developed in the early years, the Parametric approach and the Monte Carlo simulation. These papers conclude that among earlier methods, no approach appeared to perform better than the others. The evaluation and categorization of models carried out in the work by McAleer, Jimenez-Martin and Perez-Amaral(2009) and Shams and Sina (2014), among others, try to determine the conditions under which certain models predict the best. Researchers compared models in periods of varying volatility-before the crisis and after the crisis (When there was no high volatility and when volatility was high, respectively). However, this confirms that some models have good predictions before the start of the crisis, but their quality reduces with increased volatility. Others are more conservative during periods of low volatility, but in the time of the crisis the number of errors made by these models is relatively low.

Bao et al.(2006), Consigli(2002) and Danielson(2002), among others, show that in stable periods, parametric models provide satisfactory results that become less satisfactory during high volatility periods. Additional studies that find evidence in favour of parametric methods are Sarma et al.(2003), who compare Historical simulation and Parametric methods, and

Danielson and Vries(2000) in a similar comparison that also includes Extreme value theory methods. Chong(2004), who uses parametric methods to estimate VaR under a Normal distribution and under a Student's t-distribution, finds a better performance under Normality. McAleer et al.(2009) showed that RiskMetrics^{TM} was the best fitted model during a crisis, while Shams and Sina(2014) recognized GARCH(1,1) and GJR-GARCH as well forecasting models. In contrast to the results obtained by McAleer et al.(2009), the level of quality of forecasts generated by the RiskMetrics^{TM} model was considered unsatisfactory by them. However, attention needs to be drawn to one difference in the samples, on which the study was conducted, i.e. the first one comes from a developed country (USA, S&P500), and the second one from a developing country (Iran, TSEM). Taylor(2020) evaluate Value at Risk using quantile skill score and the conditional autoregressive model outperformed others.

Attempts have been made to predicts VaR with ANN. VaR estimation on the exchange rate market in the context of ANNs is dealt with in Locarek-Junge and Prinzler (1999), who illustrate how VaR estimates can be obtained by using a USD-portfolio. The empirical outcomes demonstrate an evident superiority of the neural network to other VaR models. Hamid and Iqbal(2004) compared volatility forecasts from neural networks with forecasts of implied volatility from S&P500 index futures options, using the Barone-Adesi and Whaley (BAW) American futures options pricing model. Forecasts from NN outperformed implied volatility forecasts. Similar results are put forth by He et al. (2018), who propose an innovative EMD-DBN type of ANN to estimate VaR on the USD against the AUD, CAD, CHF and the EUR. The authors find positive performance improvement in the risk estimates, and argue that the utilization of an EMD-DBN network can identify more optimal ensemble weights and is less sensitive to noise disruption compared to a FNN. Nevertheless, it is worthwhile to mention that although foreign exchange volatility forecasting through ANNs have gained some attention in the academic field, it still remains a fairly undeveloped area.

All in all, there is no full approval in the evaluation of which models should be used during periods of calm (low volatility), and which ones during crisis (High volatility).

1.3 Thesis Structure

The next chapter of discusses the properties and basic methods to estimate VaR. Subsequent chapters discuss use of Neural Network in Estimating Value at Risk and numerical comparison of the methods with examples. Findings are summarized in the last chapter.

Chapter 2

Value-at-Risk: Concept, properties and methods

2.1 Concept

Higher volatility in exchange markets, credit defaults, even endangering countries, and the call for more regulation drastically changed the circumstances in which banks operate. These situations of uncertainty are called risks and managing them is of great importance to financial institutions (e.g Banks) in order to keep them afloat. A possible method of measurement is the evaluation of losses likely to be incurred when the price of the portfolio falls. Value at Risk (VaR) does this.

According to Jorion (2001), "VaR measure is defined as the worst expected loss over a given horizon under normal market conditions at a given level of confidence. For instance, a bank might say that the daily VaR of its trading portfolio is \$2 million at the 99% confidence level. In other words, under normal market conditions, only 1% of the time, the daily loss will exceed \$2 million (99% of the time, their loss will not be more than \$2 million)". As represented in the mathematical representation below, it can also be stated as the least expected return of a portfolio at time t and at a certain level of significance, α.

Mathematically,

Let $r_1, r_2, ..., r_n$ be independently and identically distributed(iid) random variables representing financial log returns. Use $F(r)$ to denote the cumulative distribution function, $F(r) = Pr(r_t < r|\Omega_{t-1})$ conditional on the information set Ω_{t-1} available at time t-1. Assume that $\{r_t\}$ follows the stochastic process;

$$r_t = \mu_t + \varepsilon_t$$
$$\varepsilon_t = \sigma_t z_t \qquad z_i \sim N(0,1) \tag{2.1}$$

where $\sigma_t^2 = E[z_t^2|\Omega_{t-1}]$ and z_t has a conditional distribution function $G(z)$, $G(z) = Pr(z_t < z|\Omega_{t-1})$. The VaR with a given probability $\alpha \ \epsilon (0,1)$, denoted by VaR(α), is defined as the α quantile of the probability distribution of financial returns:
$F(\text{VaR}(\alpha)) = Pr(r_t < \text{VaR}(\alpha)) = \alpha$ or $\text{VaR}(\alpha) = \inf\{v|P(r_t \leq v) = \alpha\}$

One can estimate this quantile in two different ways: (1) inverting the distribution function of financial returns, F(r), and (2) inverting the distribution function of innovations, with regard to $G(z)$ the latter, it is also necessary to estimate σ_t^2.

$$\text{VaR}(\alpha) = F^{-1}(\alpha) = \mu + \sigma_t G^{-1}(\alpha) \tag{2.2}$$

Hence, a VaR model involves the specification of $F(r)$ or $G(r)$. There are several method for these estimations. Having explained the concept of Value at Risk, it is however necessary to state some of its properties or attributes.

2.2 Properties

A functional $\tau : X, Y \rightarrow \mathbb{R} \cup \{+\infty\}$ is said to be coherent risk measure for portfolios X and Y if it satisfies the following properties:

- Normalization
  $\tau[0] = 0$
  The risk when holding no assets is zero.

- Monotonicity
  if $X \leq Y$ then $\tau(X) \geq \tau(Y)$
  For financial applications, this implies that a security that always has higher return in all future states has less risk of loss.

- Translation invariance
  $\tau(X + c) = \tau(X) - c$
  In effect, if an amount of cash $\varkappa$ (or risk free asset) is added to a portfolio, then the risk is reduced by that amount.

- Positive Homogeneity
  $\tau(cX) = c\tau(X)$ if $c > 0$.
  In effect, if a portfolio or capital asset is, say, doubled, then the risk will also be doubled.

- subadditive:
  $\tau(X + Y) \leq \tau(x) + \tau(Y)$. Indeed, the risk of two portfolios together cannot get any worse than adding the two risks separately: this is the diversification principle.

One disadvantage of value at risk is that it does not always satisfy the sub-additivity, hence it might discourage diversification.

2.3 Popular methods for estimating VaR

The estimation of these functions ($F(r)$ or $G(r)$) can be carried out using the following methods:

2.3.1 Historical simulation

The historical simulation involves using past data to predict future. First of all, we have to identify the market variables that will affect the portfolio. Then, the data will be collected on the movements in these market variables over a certain time period. This provides us the alternative scenarios for what can happen between today and tomorrow. For each scenario, we calculate the changes in the dollar value of portfolio between today and tomorrow. This defines a probability distribution for changes in the value of portfolio. For instance, VaR for a portfolio using 1-day time horizon with 99% confidence level for 500 days data is nothing but an estimation of the loss when we are at the fifth-worst daily change.

Basically, historical simulation is extremely different from other type of simulation in that estimation of a covariance matrix is avoided. Therefore, this approach has simplified the computations especially for the cases of complicated portfolio.

The core of this approach is the time series of the aggregate portfolio return. More importantly, this approach can account for fat tails and is not prone to the accuracy of the

model due to being independent of model risk. As this method is very powerful and intuitive, it is then become the most widely used methods to compute VaR. However, Historical simulation requires data on all risk factors to be available over a reasonably long historical period in order to give a good representation of what might happen in the future. As it depends on history, if we run a Historical Simulations VaR in a bull market, VaR may be underestimated. Similarly, if we run a Historical Simulations VaR just after a crash, the falling returns which the portfolio has experienced recently may distort VaR.

2.3.2 GARCH Model

The Generalized Autoregressive Conditional Heteroskedasticity(GARCH) model, proposed by Bollerslev (1986) is a generalization of the ARCH process created by Engle (1982), in which the conditional variance is not only the function of lagged random errors, but also of lagged conditional variances. The standard GARCH model (p, q) can be written as:

$$r_t = \mu_t + \varepsilon_t$$
$$\varepsilon_t = \sigma_t \xi_t$$

(2.3)

where r_t = rate of return of the asset in the period t,
μ_t = conditional mean

ε_t = random error in the period t, which equals to the product of conditional standard deviation σ_t and the standardized random error ξ_t in the period t ($\xi_t \sim$ iid(0,1))

In turn, the equation of conditional variance, in the GARCH(p,q) model can be written as:

$$\sigma_t^2 = \omega + \sum_{i=1}^{q} \alpha_i \varepsilon_{t-i}^2 + \sum_{i=1}^{p} \beta_i \sigma_{t-i}^2$$

(2.4)

where σ =conditional variance in the period t,
ω = constant ($\omega > 0$)
α_i = weight of the random squared error in the period $t - 1$,
β_i = weight of the conditional variance in the period $t - 1$,
ε_{t-i}^2= squared random error in the period $t - 1$,
σ_{t-i}^2 =variance in the period $t - 1$,
q = number of random error squares periods used in the functional form of conditional variance,
p = number of lagged conditional variances used in the functional form of conditional variance.

When we use high frequency data in conjunction with GARCH models, these need to be modified to incorporate the financial market micro structure. For example, we need to incorporate heterogeneous characteristics that appear when there are many traders working in a financial market trading with different time horizons. The HARCH(n) model was introduced by Müller et al. (1997) to try to solve this problem.

2.3.3 HAR Method

High frequency data are those measured in small time intervals. This kind of data is important to study the micro structure of financial markets and also because their use is becoming feasible due to the increase of computational power and data storage. The HARCH(n) model was introduced by Müller et al. (1997) to estimate the VaR for this kind of data. In fact, this model incorporates heterogeneous characteristics of high frequency financial time series and it is given

by

$$r_t = \sigma_t \varepsilon_t$$

$$\sigma^2 = c_0 + \sum_{j=1}^{n} c_j \left(\sum_{i=1}^{j} r_{t-i} \right)^2 \qquad (2.5)$$

where $c_0 > 0, c_n > 0, c_j \geq 0 \ \forall j = 1, ..., n-1$ and ε_t are identically and independent distributed (i.i.d.) random variables with zero expectation and unit variance and the c_j are parameters estimated using least squares.

Intraday data have been found to be useful in estimating features of the distribution of daily returns. For example, the realized volatility has been used widely as a basis for forecasting the daily volatility. The heterogeneous autoregressive (HAR) model of the realized volatility is a simple and pragmatic approach, where a volatility forecast is constructed from the realized volatility over different time horizons (Corsi, 2009). However, intraday data can be expensive, and resources are required for pre-processing. Given the ready availability of the daily high and low prices, an alternative way of capturing the intraday volatility is to use the intraday range. Where Range_t is the difference between the highest and lowest log prices on day t, to predict tomorrow's range from past daily, weekly, monthly averages of Range_t, we set up the linear regression model;

$$\text{Range}_t = \beta_1 + \beta_2 \text{Range}_{t-1} + \beta_3 \text{Range}_{t-1}^w + \beta_4 \text{Range}_{t-1}^m + \varepsilon_t$$

$$\text{Range}_{t-1}^w = \frac{1}{5} \sum_{i=1}^{5} \text{Range}_{t-i} \qquad (2.6)$$

$$\text{Range}_{t-1}^m = \frac{1}{22} \sum_{i=1}^{22} \text{Range}_{t-i}$$

where Range_t is the difference between the highest and lowest log prices on day t; Range_{t-1}^w and Range_{t-1}^m are averages of Range_t over a week and month, respectively; ε_t is an i.i.d. error term with zero mean; and the β_i are parameters that are estimated using least squares. The conditional variance is then expressed as a linear function of the square of Range_t , where the intercept and the coefficient are estimated using maximum likelihoods based on a Student t distribution.

2.3.4 CaViaR Method

Engle and Manganelli (2004) propose a conditional autoregressive quantile specification (CAViaR) quantile estimation. Instead of modeling the whole distribution, the quantile is modelled directly. The empirical fact that volatilities of stock market returns cluster over time may be translated in statistical words by saying that their distribution is autocorrelated. Consequently, the VaR, which is a quantile, must behave in similar way. A natural way to formalize this characteristic is to use some type of autoregressive specification

let $\{y_t\}_{t=1}^T$ and θ be a vector of portfolio returns and the probability associated with VaR respectively. Let x_t be a vector of time t observable variables(return or any other observable variables), and β_θ be a p-vector of unknown parameters. Finally, let $f_t(\beta) \equiv f_t(x_{t-1}, \beta_\theta)$ denote the time t θ-quantile of the distribution of the portfolio returns formed at $t-1$, The θ subscript is however suppressed from β_θ for convenience in notation: The general

specification of CaViar would be:

$$f_t(\beta) = \beta_0 + \sum_{i=1}^{q} \beta_i f_{t-i}(\beta) + \sum_{j=1}^{r} \beta_j l(x_{t-j}) \tag{2.7}$$

where $p = q + r + 1$ is the dimension of β and l is a function of a finite number of lagged values of observables. The autoregressive terms $\beta_i f_{t-i}(\beta)$, $i = 1, ..., q$, ensure that the quantile changes smoothly over time. The parameters of CaViaR are estimated by quantile regression. The role of $l(x_{t-j})$ is to connect $f_t(\beta)$ to observable variables that belong to the information set.

The asymmetric slope is a variant of CaViar model, which allows the response to positive and negative returns to be different. It is modelled as:

$f_t(\beta) = \beta_1 + \beta_2 f_{t-1}(\beta) + \beta_3 (y_{t-1})^+ + \beta_4 (y_{t-1})^-$.   where, $(y_{t-1})^+ = \max(y_{t-1},\ 0)$ and $(y_{t-1})^- = \min(y_{t-1}, 0)$

Note: In this work, value at risk will be estimated based on financial log-returns.

Chapter 3

Estimating VaR using Neural Networks

Neural networks, also known as artificial neural networks (ANNs) are computing systems vaguely inspired by the biological neural networks that constitute animal brains. Their name and structure are inspired by the human brain, mimicking the way that biological neurons signal to one another.

Neural networks are made of node layers, containing an input layer, one or more hidden layers, and an output layer. Each node (or artificial neuron), connects to another and has an assigned weight and threshold. If the output of any individual node exceeds the specified threshold value, that node is activated, transferring data to the next layer of the network. Else, no data is passed along to the next layer of the network.


Figure 3.1: A figure showing the layers of a Neural Network (see IBM)

Neural networks rely on training data to learn and improve their accuracy over time. However, once these learning algorithms are polished for accuracy, they become powerful tools in computer science and artificial intelligence, allowing us to classify and cluster data at a high speed. Tasks in speech recognition or image recognition can take minutes versus hours when compared to the manual identification by human experts.

One major advantage of Neural Networks is that they have a flexible nonlinear function mapping capability, meaning that any model can be arbitrarily well approximated by a

sufficiently large ANN. This model-free property of ANNs is an important advantage since returns do not display a specific nonlinear pattern. However, just like historical simulation they require large data set (which is not always available) for training to perform excellently well.

3.1 Mathematics of Neural Network

Here, we consider a case of single neuron

3.1.1 A single Neuron

Figure 3.2: A single neuron of neural networks

Figure 3.2 shows a network with one layer containing a single neuron. This neuron receives input from the prior input layer, performs computations, and gives output. x_1 and x_2 are inputs with weights w_1 and w_2 respectively. The neuron applies a function f to the dot-product of these inputs, which is $w_1 x1 + w_2 x_2 + b$. Besides the two numerical input values, there is one input value 1 with weight b, called the Bias. The main function of bias is to stand for unknown parameters or unforeseen factors. The output Y is computed by taking the dot-product of all input values and their associated weights and putting it into the function f. This function is called the Activation Function.

Activation functions are needed because many problems take multiple influencing factors into account and yield classifications. For example, if one encounters a binary classification problem, the results would be either yes or no, activation functions are needed to map the results inside this range. If one encounters a problem involving probability, then one would wish to see the predictions from the neural network being in the range of $[0, 1]$. This is what activation functions can do.

There are two types of activation functions: linear activation functions and non-linear ones. The biggest limitation of linear ones is that they cannot learn complex function mappings because they are just polynomials of one degree. Therefore, we always need non-linear activation functions to produce results in desired ranges and to send them as inputs to the next layer. The following subsection will introduce few generally used non-linear activation functions.

3.1.2 Activation Functions

An activation function takes the dot-product mentioned before as a input and performs a certain computation on it. We put a certain activation function inside of neurons of hidden layers based on the range of the result we expect to see. A notable property of activation

functions is that they should be differentiable, because later we need this property to train the neural network using backpropagation optimization.

Here are few commonly used activation functions:

Sigmoid: This takes a real-valued input and returns a output in the range [0,1]:

$\delta = \frac{1}{1+e^{-x}}$

Figure 3.3: Sigmoid() Activation Function

Figure 3.3 shows an S-shaped curve and the values going through the Sigmoid function will be squeezed in the range of [0, 1]. Since the probability of anything exists only between the range of 0 and 1, Sigmoid is a compatible transfer function for probability. Although the Sigmoid function is easy to understand and ready to use, it is not frequently used because it has vanishing gradient problem. This problem is that, in some cases, the gradient gets so close to zero that it does not effectively apply change to the weight. In the worst case, this may completely stop the neural network from further training. Second, the output of this function is not zero-centered, which makes the gradient updates go far in different directions. Besides, the fact that output is in the narrow range [0, 1] makes optimization harder. In order to compensate the shortcomings, tanh() is an alternative option because it is a stretched version of the Sigmoid function, in which its output is zero-centered.

tanh: This takes real-valued input and produces the results in the range [-1, 1]:

$\tanh(x) = \frac{sinh(x)}{cosh(x)} = \frac{e^x - e^{-x}}{e^x + e^{-x}}$

Figure 3.4: tanh() Activation Function

The advantage is that the negative input values will be mapped strongly negative and the

zeros will be mapped near zero through this function. Therefore, this function is useful in the performance of a classification between two distinct classes. Though, this function is preferred over the Sigmoid function in practice (output in wider range), but the gradient vanishing problem still exists. The following ReLU function rectifies this problem using a relatively simple formula.

ReLU (Rectified Linear Unit): It takes a real-valued input and replaces the negative values with zero:

$R(x) = \max(0, x)$

Figure 3.5: ReLU() Activation Function

It is used in almost all the convolutional neural networks or deep learning because it is a relatively simple and efficient function which avoids and rectifies the gradient vanishing problem. The problem of this activation function is that all the negative values become zeros after this activation, which in turns affects the results by not taking negative values into account.

We use different activation functions when we know what characteristics of results we expect to see.

3.2 General Model Building

Having discussed the mathematics behind neural network, it is however in important to talk about the neural network architectures and other components

3.2.1 Neural network Architures

Neural Networks are complex structures made of artificial neurons that can take in several inputs to produce a single(or more) output(s). Normally, a Neural Network consists of an input and output layer with one or multiple hidden layers within. In a Neural Network, all the neurons(contained in each layers) influence each other, and hence, they are all connected. The manner in which the input neurons produce a certain output is intimately linked to the structure of the neural network. The two main classes of network architectures are discussed below

Feed-forward Neural Network

In feed-forward neural network(FFN), each neuron in a particular layer is connected with all neurons in a subsequent layer. The information flow in the network is of feedforward type (i.e the connections can never skip a layer, or form any loops backwards). The picture below shows

Figure 3.6: A fully connected FFN with a single hidden layer (see Anna & Bijelic, 2019)

the architectural layout of a single hidden layer feed-forward neural network.

As shown the above figure, the values of the input are transported to the hidden layer through connections, each being characterised by certain weight coefficient, $W_{i,k}$. The degree of connection between an the input node and a hidden node is reflected by these weight coefficients. Defining $[x_{1,t}; x_{2,t}; ...; x_{n,t}]$ as the vector of the input signals and $[h_{1,t}; h_{2,t}; ...; h_{m,t}]$, the propagation of the input nodes to one hidden node can mathematically be described by:

$$h_{k,t} = \sum_{i=1}^{n} W_{i,k} \cdot x_{i,t} \text{ for } k = 1, 2, ..., m$$

An undesirable property of the formula is its linear representation, which, if applied, would suggest that the output prediction would be a linear function, which is not always the case. In order to deal with this, a non-linear activation function, $\Phi(\cdot)$ is applied to the weighted sum of inputs into a hidden node. This activation function, which in the majority of applications takes the form of a sigmoid function or a ReLu function, makes the neural network capable of approximating virtually any function. However, before applying the activation function, a bias vector $[b_1; b_2; ...; b_m]$ is added, which essentially indicates whether a neuron tends to be active or inactive in the prediction process. The propagation from the input layer to the hidden layer in a feed-forward neural network can then be reformulated to:

$$h_{k,t} = \Phi(b_{k,0} + \sum_{i=1}^{n} W_{i,k} \cdot x_{i,t}) \text{ for } k = 1, 2, ..., m$$

Although the model-free assumption underlying the feedforward neural network theoretically suggests that it should far better than the conventional GARCH(1,1) in predicting volatility, the feedforward neural network is subject to a major concern in that it precludes modelling time-dependencies in the data. This deficiency of not being able to take correlations between inputs into account is however resolved in the recurrent neural network, which is able to selectively pass information across sequences of elements by creating cycles in the network.

Recurrent Neural Network

Recurrent Neural Networks (RNN) can handle sequential data due to the capability of each neuron to maintain information about previous inputs, contrary feedforward neural networks. This means that the prediction a recurrent neural network node made at previous time step
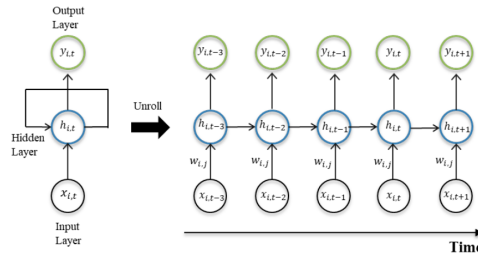
Figure 3.7: Representation of an unrolled plain vanilla recurrent neural network. (see Anna & Bijelic, 2019)

$t - 1$ affects the prediction it will make one moment later, at time step t. RNN nodes can be thought of as having memory as it takes inputs not only the current signal, but also what has been perceived previously in time.

RNNs contain feedback loops from the so-called hidden states, and this allows preservation of information from one node to another while reading in inputs. This feedback loop mechanism occurs at each time step in the data series, which results in each hidden state containing traces not only of the previous hidden state, but also of all the preceding ones, for as long as the memory of the network persists (Skymind, 2019).
The unrolled RNN illustrates how the network allows the hidden neurons to see their own previous output, so that their subsequent behavior can be shaped by previous responses (Tenti, 1996). Furthermore, by introducing time-lagged model components, it becomes evident that the utilization of a RNN is particularly desired when there are time dependencies in the data series at hand. Using the previous notation and assuming that the hidden states are the ones looped back, the output from a hidden node in the RNN model depends on the input values at time t, but also on its own lagged values at order p as shown below:

$$h_{k,t} = \Phi(b_{k,0} + \sum_{i=1}^{n} W_{i,k} \cdot x_{i,t}) + \sum_{k=1}^{m} \gamma_k \cdot h_{k,t-p} \text{for } k = 1, 2, ..., m$$

where $h_{k,t-p}$ represents the lagged hidden state values at order p, and γ_k a coefficient. Another distinguishing characteristic of recurrent networks is that they share parameters across each layer of the network. While feedforward networks have different weights across each node, recurrent neural networks share the same weight parameter within each layer of the network. That said, these weights are still adjusted in the through the processes of backpropagation and gradient descent to facilitate reinforcement learning. Recurrent neural networks leverage backpropagation through time (BPTT) algorithm to determine the gradients, which is slightly different from traditional backpropagation as it is specific to sequence data. The principles of BPTT are the same as traditional backpropagation, where the model trains itself by calculating errors from its output layer to its input layer. These calculations allow us to adjust and fit the parameters of the model appropriately. BPTT differs from the traditional approach in that BPTT sums errors at each time step whereas feedforward networks do not need to sum errors as they do not share parameters across each layer.

Through this process, RNNs tend to run into two problems, known as exploding gradients and vanishing gradients. These issues are defined by the size of the gradient, which is the slope of the loss function along the error curve. When the gradient is too small, it continues to become smaller, updating the weight parameters until they become insignificant—i.e. 0. When that occurs, the algorithm is no longer learning. Exploding gradients occur when the gradient is too large, creating an unstable model. In this case, the model weights will grow too large, and they will eventually be represented as NaN. One solution to these issues is to reduce the number of hidden layers within the neural network, eliminating some of the complexity in the RNN model(see, IBM 2020).

Long Short-Term Memory Recurrent Neural Network

Originally introduced by Hochreiter and Schmidhuber (1997), a main characteristic of LSTMs
– which are a sub class of recurrent neural networks - is its purpose-built memory cells, which
allows it to capture long range dependencies in the data. From a model perspective, LSTMs
differ from other neural network architectures in that they are applied recurrently

The output from a previous sequence of the network function serves – in combination
with the next sequence element - as input for the next application of the network function. In
this sense, the LSTM can be interpreted as being similar to an HMM (Hidden Markov Model),
in that there is a hidden state which conditions the output distribution. However, the LSTM
hidden state not only depends on its previous states, but it also captures long term sequence
dependencies through its recurrent nature. Maybe most notably, the receptive field size of an
LSTM is not bound architecture wise as in case of simple feed forward network and CNN.
Instead, the LSTM's receptive field depends solely on the LSTMs ability to memorize the past
input.

Chapter 4

Figures, tables, enumerate and itemize

4.1 Figures

In almost every document figures will be needed in order to explain a concept or just present something. The package *graphicx* is needed for embedding figures.


Figure 4.1: A figure caption beneath the figure for description of the depicted concept which sometimes can be very long

In Figure 4.1, for example, a PNG image is depicted (compiled with pdflatex). Alternatively, EPS figures can be embedded if dvips and ps2pdf compilation is used. All figures are listed in the list of figures with the command *listoffigures*.

4.2 Tables

Data can be presented in tables, e.g., as shown in Table 4.1.

	Property 1	Property 2
Criterion 1	764	23546
Criterion 2	3	34

Table 4.1: Exemplary table

Sometimes very long tables must be presented which may also go over pages. For this cases the packages *longtable* is useful, as used in

i^3	$2i^3$	$3i^3$
1	2	3
8	16	24
27	54	81
64	128	192
125	250	375
216	432	648
343	686	1029
512	1024	1536
729	1458	2187
1000	2000	3000
1331	2662	3993
1728	3456	5184
2197	4394	6591
2744	5488	8232
3375	6750	10125
4096	8192	12288
4913	9826	14739
5832	11664	17496
6859	13718	20577
8000	16000	24000
9261	18522	27783
10648	21296	31944
12167	24334	36501
13824	27648	41472
15625	31250	46875
17576	35152	52728
19683	39366	59049
21952	43904	65856
24389	48778	73167
27000	54000	81000
29791	59582	89373
32768	65536	98304
35937	71874	107811
39304	78608	117912
42875	85750	128625
46656	93312	139968
50653	101306	151959
54872	109744	164616
59319	118638	177957
64000	128000	192000
68921	137842	206763
74088	148176	222264
79507	159014	238521
85184	170368	255552
91125	182250	273375
97336	194672	292008
103823	207646	311469
110592	221184	331776
117649	235298	352947
125000	250000	375000

Table 4.2: Long Table

All tables are listed with *listoftables*.

4.3   Enumerate and itemize

If important sequential points are to presented the environment *enumerate* can be used as follows:

1. Some important stuff

2. More stuff

With the package *enumerate* some options can be used, e.g.,

a) Some important stuff

b) More stuff

or

1) Some important stuff

2) More stuff

Alternatively, point can be just presented without any enumeration with the environment *itemize*

- Some important stuff

- More stuff

Chapter 5

Appendix, footnotes, todos and index

5.1 Appendix

For many reasons some concept may be important for the document but too long for the main text. In this kind of cases these concept can be presented with the environment *appendix* in appendices, e.g., as in Appendix A and Appendix B.

5.2 Footnotes

You may want to give additional information to some points[1] in the text[2].

5.3 Todos

With the package *todonotes* comments pointing to their place can be embedded into the text. These comments are veeeery useful if you are writing something for the first time or are working on a draft. The todos can be listed with *listoftodos* where you want it to appear in order to see what is unfinished or needs some more work.

like this one

5.4 Index

If the document is very long, it may be very useful for a lot of readers to have an index for searching key words and certain concepts (Crtl+F is usually very helpful in PDFs but not always the best solution). For this, the package *makeidx*, the commands *makeindex* and *printindex* and the compiling option *make index* are needed. You may want to index different words like heterogeneous materials, effective properties and homogenization.

---

[1]Bla bla

[2]Blu blup

Appendix A

Just an example appendix

A.1 Bla blup

Sme stuff

$$f(x) = \int_\Omega g(x)dx \ .$$ \hfill (A.1)

Appendix B

Another example

B.1   More stuff

Bla bla.

Bibliography

R. Hill. The elastic behaviour of a crystalline aggregate. *Proceedings of the Physical Society. Section A*, 65:349–354, 1952.

E. Kröner. Bounds for effective elastic moduli of disordered materials. *Journal of the Mechanics and Physics of Solids*, 25(3):137–155, 1977.

List of Figures

List of Tables

Index